

Super resolution

```
In [ ]: 1 import fastai
        2 from fastai.vision import *
        3 from fastai.callbacks import *
        4
        5 from torchvision.models import vgg16_bn
```

```
In [ ]: 1 path = untar_data(URLs.PETS)
        2 path_hr = path/'images'
        3 path_lr = path/'small-96'
        4 path_mr = path/'small-256'
```

```
In [ ]: 1 il = ImageItemList.from_folder(path_hr)
```

```
In [ ]: 1 def resize_one(fn,i):
        2     dest = path_lr/fn.relative_to(path_hr)
        3     dest.parent.mkdir(parents=True, exist_ok=True)
        4     img = PIL.Image.open(fn)
        5     targ_sz = resize_to(img, 96, use_min=True)
        6     img = img.resize(targ_sz, resample=PIL.Image.BILINEAR).convert('RGB')
        7     img.save(dest, quality=60)
```




```
In [ ]: 1 # to create smaller images, uncomment the next line when you run this the fi
        2 # parallel(resize_one, il.items)
```

```
In [ ]: 1 bs,size=32,128
        2 arch = models.resnet34
        3
        4 src = ImageImageList.from_folder(path_lr).random_split_by_pct(0.1, seed=42)
```

```
In [ ]: 1 def get_data(bs,size):
        2     data = (src.label_from_func(lambda x: path_hr/x.name)
        3             .transform(get_transforms(max_zoom=2.), size=size, tfm_y=True)
        4             .databunch(bs=bs).normalize(imagenet_stats, do_y=True))
        5
        6     data.c = 3
        7     return data
```

```
In [ ]: 1 data = get_data(bs,size)
```

```
In [ ]: 1 data.show_batch(ds_type=DatasetType.Valid, rows=2, figsize=(9,9))
```



Feature loss

```
In [ ]: 1 t = data.valid_ds[0][1].data  
2 t = torch.stack([t,t])
```

```
In [ ]: 1 def gram_matrix(x):  
2     n,c,h,w = x.size()  
3     x = x.view(n, c, -1)  
4     return (x @ x.transpose(1,2))/(c*h*w)
```

```
In [ ]: 1 gram_matrix(t)
```

```
Out[ ]: tensor([[[[0.0759, 0.0711, 0.0643],
                  [0.0711, 0.0672, 0.0614],
                  [0.0643, 0.0614, 0.0573]],

                [[0.0759, 0.0711, 0.0643],
                  [0.0711, 0.0672, 0.0614],
                  [0.0643, 0.0614, 0.0573]]]])
```

```
In [ ]: 1 base_loss = F.l1_loss
```

```
In [ ]: 1 vgg_m = vgg16_bn(True).features.cuda().eval()
2 requires_grad(vgg_m, False)
```

```
In [ ]: 1 blocks = [i-1 for i,o in enumerate(children(vgg_m)) if isinstance(o,nn.MaxPo
2 blocks, [vgg_m[i] for i in blocks]
```

```
Out[ ]: ([5, 12, 22, 32, 42],
         [ReLU(inplace), ReLU(inplace), ReLU(inplace), ReLU(inplace), ReLU(inplace)])
```

```
In [ ]: 1 class FeatureLoss(nn.Module):
2     def __init__(self, m_feat, layer_ids, layer_wgts):
3         super().__init__()
4         self.m_feat = m_feat
5         self.loss_features = [self.m_feat[i] for i in layer_ids]
6         self.hooks = hook_outputs(self.loss_features, detach=False)
7         self.wgts = layer_wgts
8         self.metric_names = ['pixel',] + [f'feat_{i}' for i in range(len(layer_ids))]
9         + [f'gram_{i}' for i in range(len(layer_ids))]
10
11     def make_features(self, x, clone=False):
12         self.m_feat(x)
13         return [(o.clone() if clone else o) for o in self.hooks.stored]
14
15     def forward(self, input, target):
16         out_feat = self.make_features(target, clone=True)
17         in_feat = self.make_features(input)
18         self.feat_losses = [base_loss(input,target)]
19         self.feat_losses += [base_loss(f_in, f_out)*w
20                               for f_in, f_out, w in zip(in_feat, out_feat, self.wgts)]
21         self.feat_losses += [base_loss(gram_matrix(f_in), gram_matrix(f_out))
22                               for f_in, f_out, w in zip(in_feat, out_feat, self.wgts)]
23         self.metrics = dict(zip(self.metric_names, self.feat_losses))
24         return sum(self.feat_losses)
25
26     def __del__(self): self.hooks.remove()
```


```
In [ ]: 1 feat_loss = FeatureLoss(vgg_m, blocks[2:5], [5,15,2])
```

Train

```
In [ ]: 1 wd = 1e-3
        2 learn = unet_learner(data, arch, wd=wd, loss_func=feat_loss, callback_fns=[Lo
        3                                     blur=True, norm_type=NormType.Weight)
        4 gc.collect();
```

```
In [ ]: 1 learn.lr_find()
        2 learn.recorder.plot()
```

LR Finder is complete, type {learner_name}.recorder.plot() to see the graph.



```
In [ ]: 1 lr = 1e-3
```




```
In [ ]: 1 def do_fit(save_name, lrs=slice(lr), pct_start=0.9):
        2     learn.fit_one_cycle(10, lrs, pct_start=pct_start)
        3     learn.save(save_name)
        4     learn.show_results(rows=1, imgsize=5)
```

```
In [ ]: 1 do_fit('1a', slice(lr*10))
```

Total time: 11:16

epoch	train_loss	valid_loss	pixel	feat_0	feat_1	feat_2	gram_0
1	3.873667	3.759143	0.144560	0.229806	0.314573	0.226204	0.552578
2	3.756051	3.650393	0.145068	0.228509	0.308807	0.218000	0.534508
3	3.688726	3.628370	0.157359	0.226753	0.304955	0.215417	0.522482
4	3.628276	3.524132	0.145285	0.225455	0.300169	0.211110	0.497361
5	3.586930	3.422895	0.145161	0.224946	0.294471	0.205117	0.472445
6	3.528042	3.394804	0.142262	0.220709	0.289961	0.201980	0.478097
7	3.522416	3.361185	0.139654	0.220379	0.288046	0.200114	0.471151
8	3.469142	3.338554	0.142112	0.219271	0.287442	0.199255	0.462878
9	3.418641	3.318710	0.146493	0.219915	0.284979	0.197340	0.455503
10	3.356641	3.187186	0.135588	0.215685	0.277398	0.189562	0.432491

Input / Prediction / Target






```
In [ ]: 1 learn.unfreeze()
```

```
In [ ]: 1 do_fit('1b', slice(1e-5,1r))
```

Total time: 11:39

epoch	train_loss	valid_loss	pixel	feat_0	feat_1	feat_2	gram_0
1	3.303951	3.179916	0.135630	0.216009	0.277359	0.189097	0.430012
2	3.308164	3.174482	0.135740	0.215970	0.277178	0.188737	0.428630
3	3.294504	3.169184	0.135216	0.215401	0.276744	0.188395	0.428544
4	3.282376	3.160698	0.134830	0.215049	0.275767	0.187716	0.427314
5	3.301212	3.168623	0.135134	0.215388	0.276196	0.188382	0.427277
6	3.299340	3.159537	0.135039	0.214692	0.275285	0.187554	0.427840
7	3.291041	3.159207	0.134602	0.214618	0.275053	0.187660	0.428083
8	3.285271	3.147745	0.134923	0.214514	0.274702	0.187147	0.423032
9	3.279353	3.138624	0.136035	0.213191	0.273899	0.186854	0.420070
10	3.261495	3.124737	0.135016	0.213681	0.273402	0.185922	0.416460

Input / Prediction / Target



```
In [ ]: 1 data = get_data(12,size*2)
```

```
In [ ]: 1 learn.data = data
        2 learn.freeze()
        3 gc.collect()
```

Out[]: 0



```
In [ ]: 1 learn.load('1b');
```

```
In [ ]: 1 do_fit('2a')
```

Total time: 43:44

epoch	train_loss	valid_loss	pixel	feat_0	feat_1	feat_2	gram_0
1	2.249253	2.214517	0.164514	0.260366	0.294164	0.155227	0.385168
2	2.205854	2.194439	0.165290	0.260485	0.293195	0.154746	0.374004
3	2.184805	2.165699	0.165945	0.260999	0.291515	0.153438	0.361207
4	2.145655	2.159977	0.167295	0.260605	0.290226	0.152415	0.359476
5	2.141847	2.134954	0.168590	0.260219	0.288206	0.151237	0.348900
6	2.145108	2.128984	0.164906	0.259023	0.286386	0.150245	0.352594
7	2.115003	2.125632	0.169696	0.259949	0.286435	0.150898	0.344849
8	2.109859	2.111335	0.166503	0.258512	0.283750	0.148191	0.347635
9	2.092685	2.097898	0.169842	0.259169	0.284757	0.148156	0.333462
10	2.061421	2.080940	0.167636	0.257998	0.282682	0.147471	0.330893

Input / Prediction / Target





```
In [ ]: 1 learn.unfreeze()
```

```
In [ ]: 1 do_fit('2b', slice(1e-6,1e-4), pct_start=0.3)
```

Total time: 45:19

epoch	train_loss	valid_loss	pixel	feat_0	feat_1	feat_2	gram_0
1	2.061799	2.078714	0.167578	0.257674	0.282523	0.147208	0.330824
2	2.063589	2.077507	0.167022	0.257501	0.282275	0.146879	0.331494
3	2.057191	2.074605	0.167656	0.257041	0.282204	0.146925	0.330117
4	2.050781	2.073395	0.166610	0.256625	0.281680	0.146585	0.331580
5	2.054705	2.068747	0.167527	0.257295	0.281612	0.146392	0.327932
6	2.052745	2.067573	0.167166	0.256741	0.281354	0.146101	0.328510
7	2.051863	2.067076	0.167222	0.257276	0.281607	0.146188	0.327575
8	2.046788	2.064326	0.167110	0.257002	0.281313	0.146055	0.326947
9	2.054460	2.065581	0.167222	0.257077	0.281246	0.146016	0.327586
10	2.052605	2.064459	0.166879	0.256835	0.281252	0.146135	0.327505

Input / Prediction / Target



Test

```
In [ ]: 1 learn = None
        2 gc.collect();
```

```
In [ ]: 1 256/320*1024
```

Out[]: 819.2

```
In [ ]: 1 256/320*1600
```

Out[]: 1280.0

```
In [ ]: 1 learn = unet_learner(data, arch, loss_func=F.l1_loss, blur=True, norm_type=N
```



```
In [ ]: 1 data_mr = (ImageImageList.from_folder(path_mr).random_split_by_pct(0.1, seed
2         .label_from_func(lambda x: path_hr/x.name)
3         .transform(get_transforms(), size=(1280,1600), tfm_y=True)
4         .databunch(bs=1).normalize(imagenet_stats, do_y=True))
5 data_mr.c = 3
```

```
In [ ]: 1 learn.load('2b');
```

```
In [ ]: 1 learn.data = data_mr
```

```
In [ ]: 1 fn = data_mr.valid_ds.x.items[0]; fn
```

```
Out[ ]: PosixPath('/data1/jhoward/git/course-v3/nbs/dl1/data/oxford-iiit-pet/small-256/
Siamese_178.jpg')
```

```
In [ ]: 1 img = open_image(fn); img.shape
```

```
Out[ ]: torch.Size([3, 256, 320])
```

```
In [ ]: 1 p,img_hr,b = learn.predict(img)
```

```
In [ ]: 1 show_image(img, figsize=(18,15), interpolation='nearest');
```



```
In [ ]: 1 Image(img_hr).show(figsize=(18,15))
```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).



```
In [ ]: 1
```